



## Lecture 13

# Rust and WebAssembly

# WebAssembly Demo

# What is wrong with JavaScript?

- The strength is a WEAKNESS:

**Easy to learn + Executes in browser + Dynamic Typing**

=

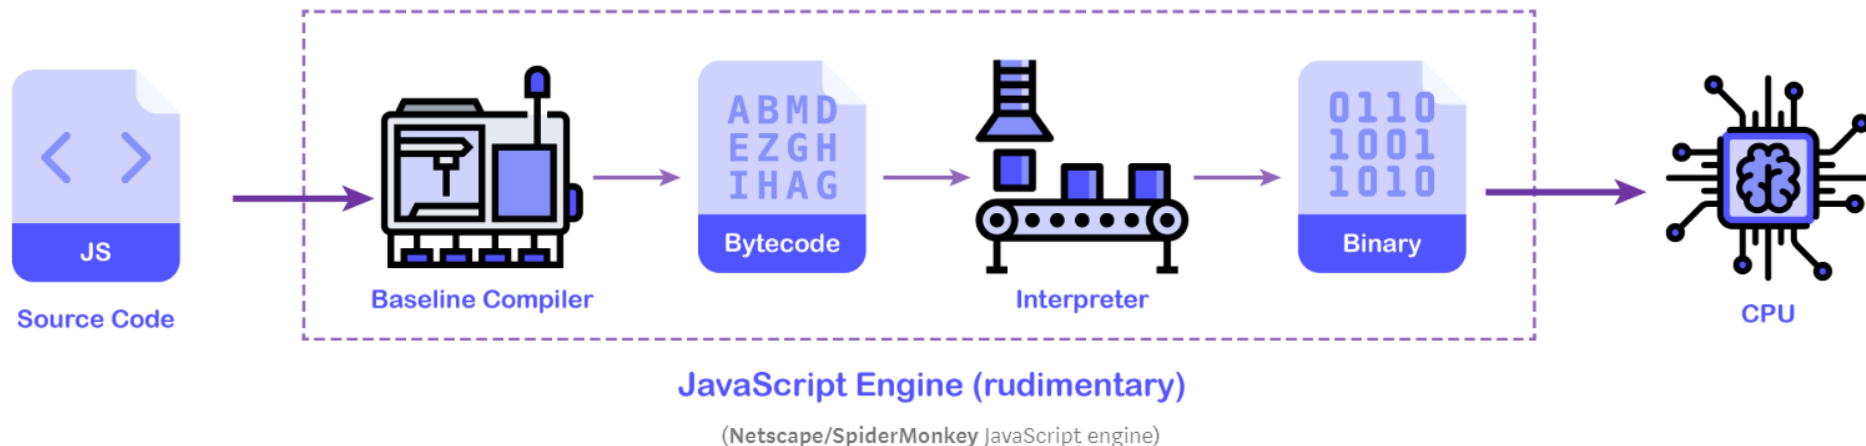
**LOW PERFORMANCE**

- 3D games? image/video editing?
- The cost of downloading, parsing, and compiling very large JavaScript applications can be prohibitive.
- Mobile and other resource-constrained platforms?

# How does JavaScript work?

- Every browser provides a JavaScript engine that runs the JavaScript code.

## *Rudimentary*

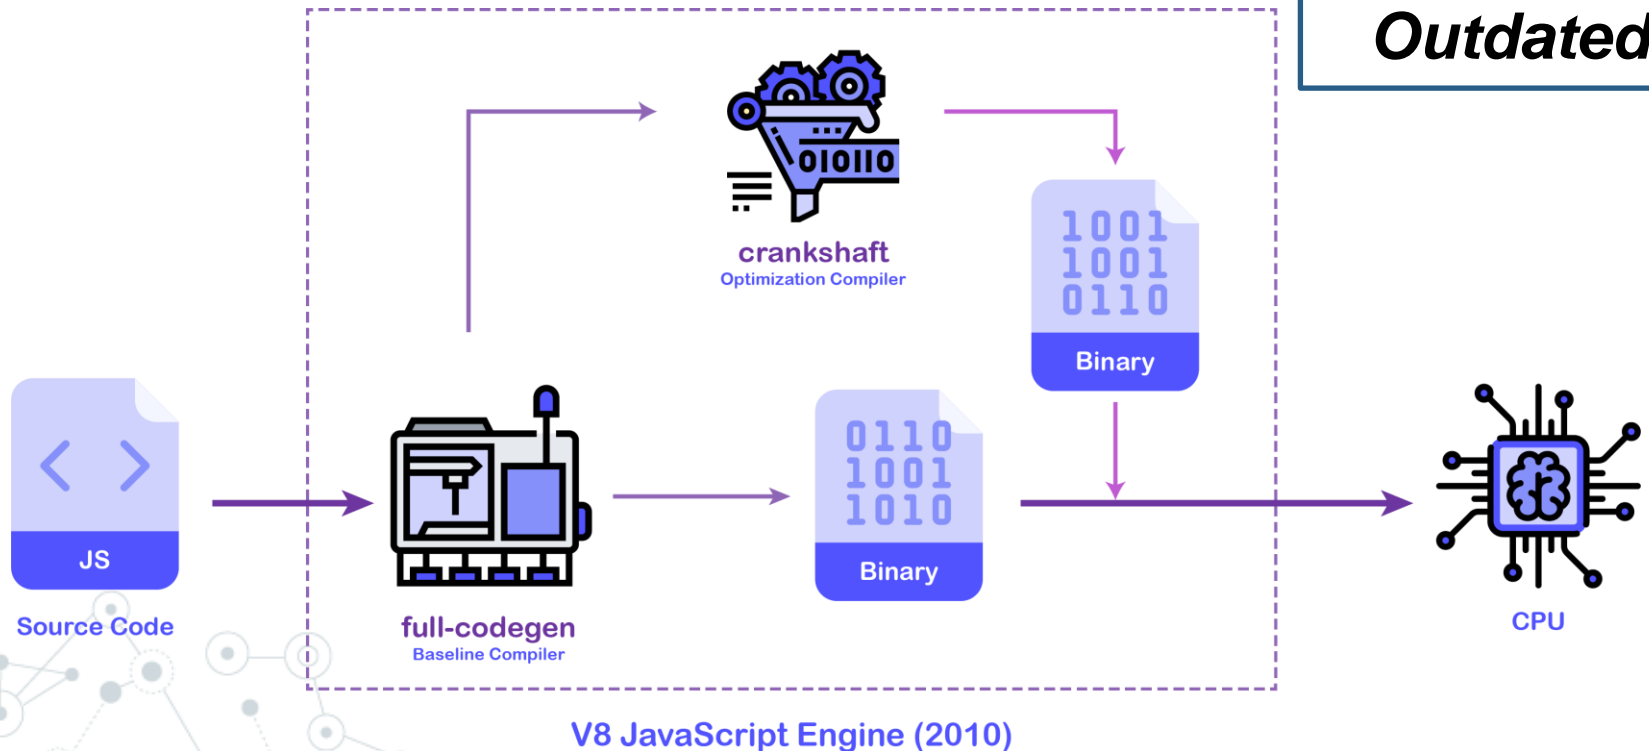


- When it comes to a highly dynamic and interactive web application, the user experience is very poor with this model.

# How does JavaScript work?

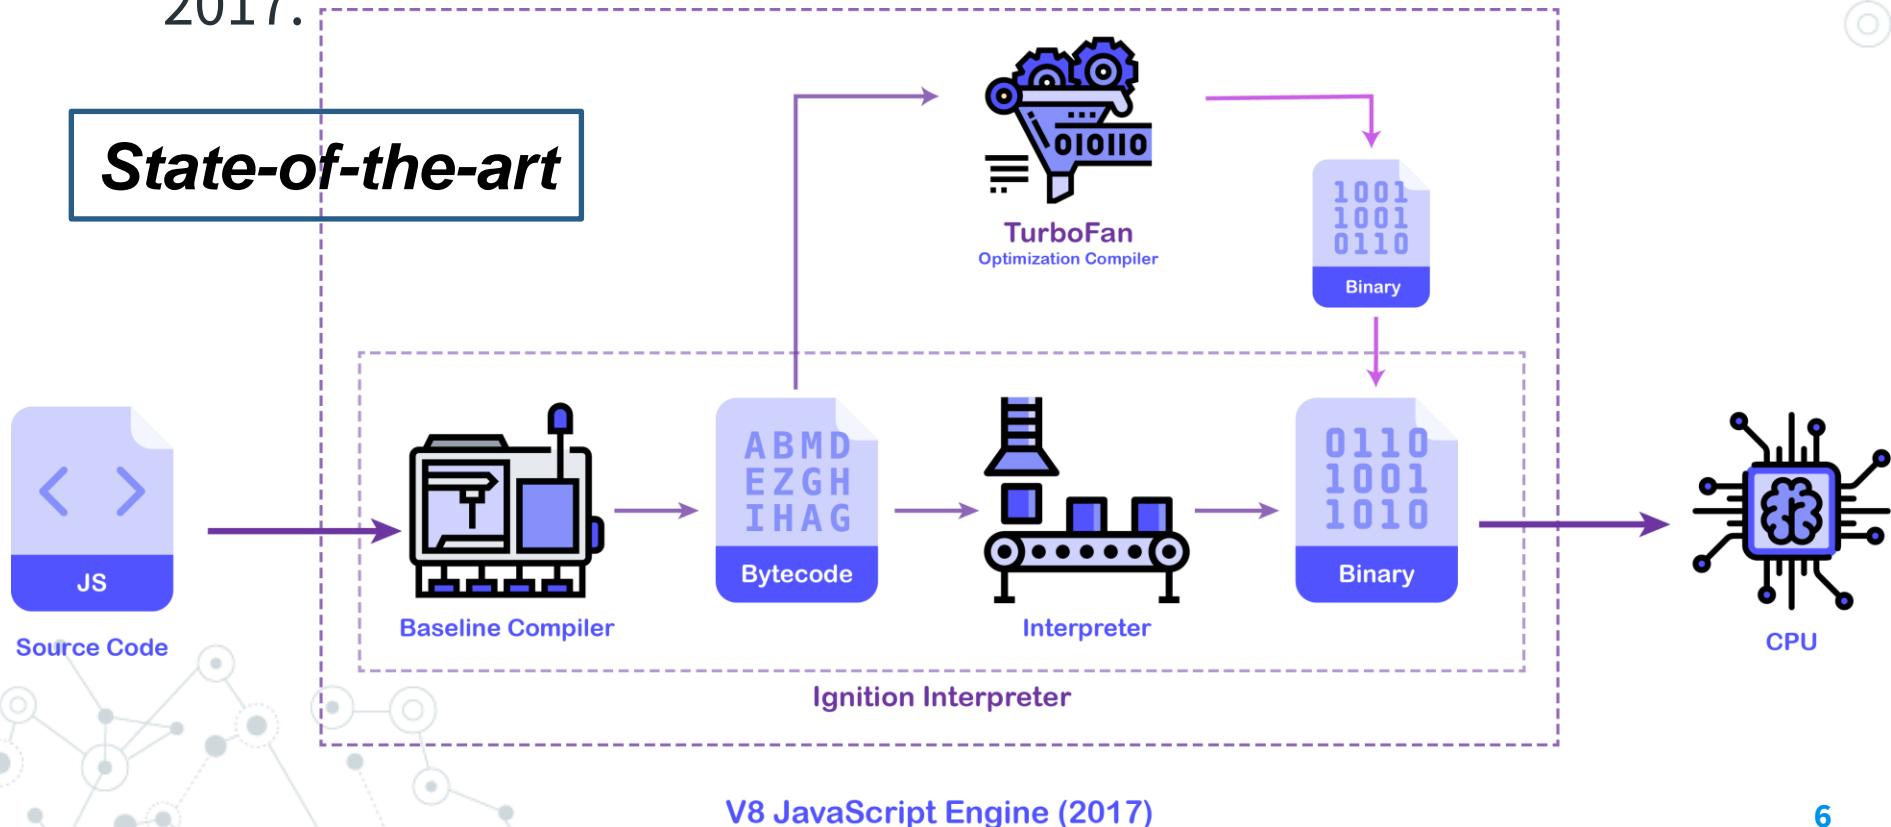
- This problem was faced by Google's Chrome browser while displaying Google Maps on the web.
- To increase the JavaScript performance on the web, they used the V8 JavaScript engine.

***Outdated***



# How JavaScript is optimized?

- The V8 team created a new version of the V8 engine from the ground up.
- This new version of the JavaScript engine was released in 2017.





# WebAssembly: Game Changer!

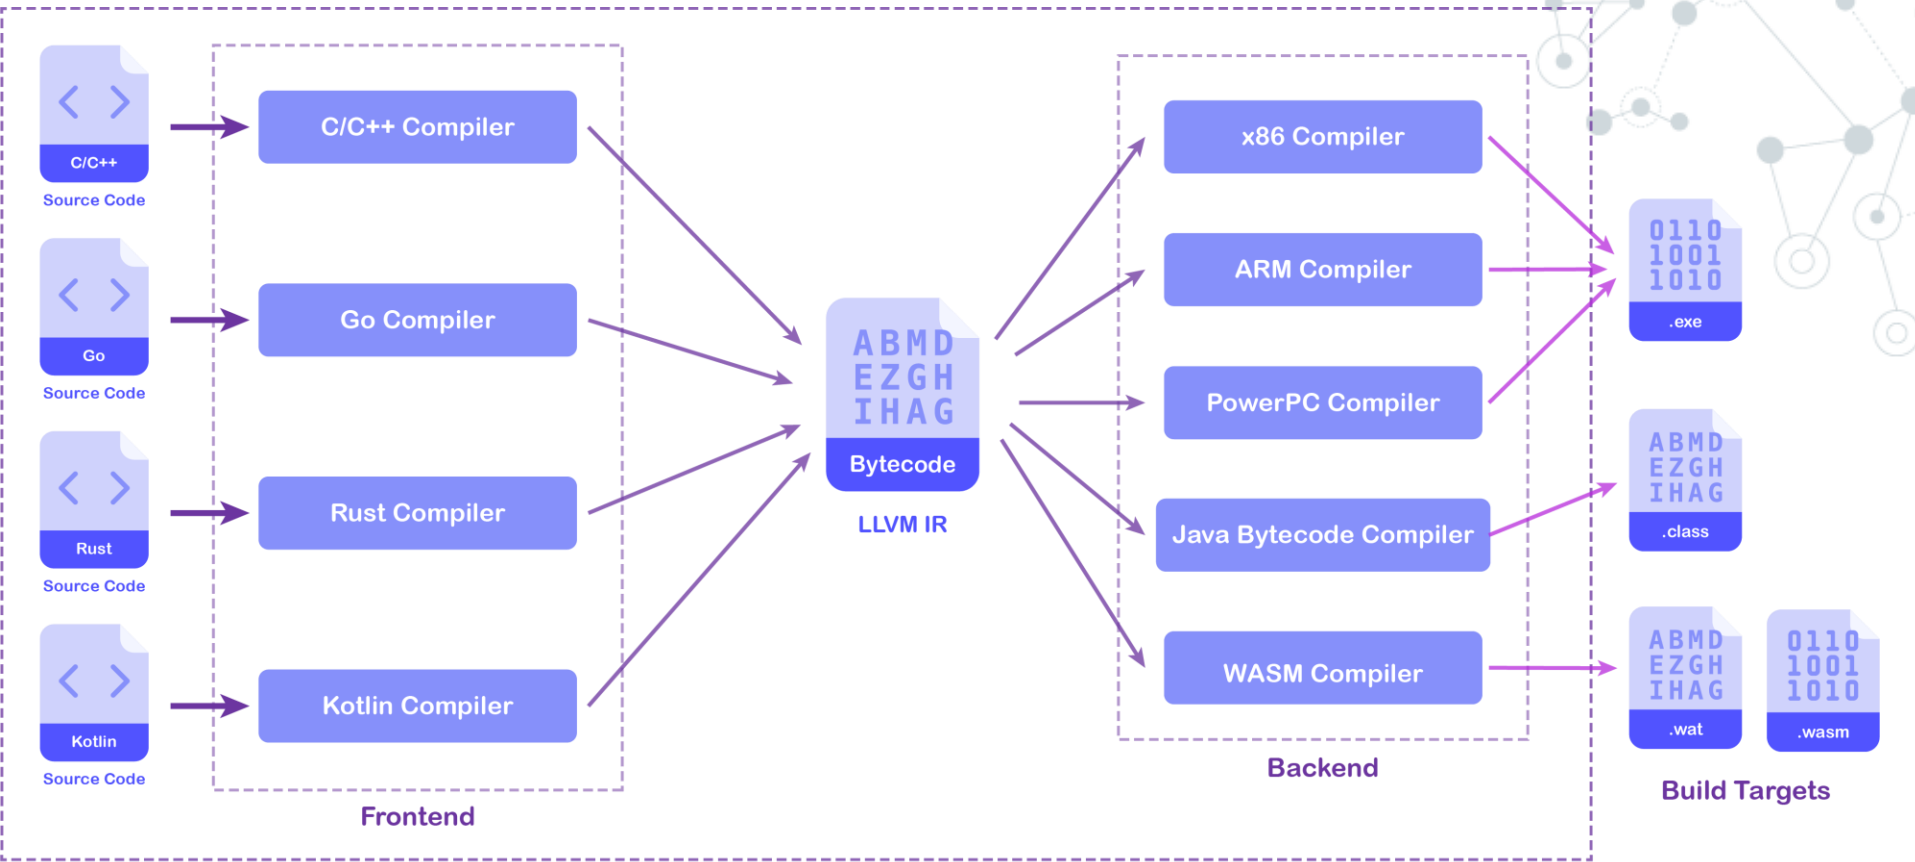
# WebAssembly: Game Changer!

- WebAssembly, a bytecode standard for web browsers.
- Announced in 17 June 2015.
- Developed by WebAssembly Working Group.
- Build target
- Binary format
- Supports JS.



# WebAssembly: Game Changer!

- A language for the web.
- Compiled from other languages.
- Offers maximized, reliable performance.
- Not a replacement for JS.
- It is meant to augment the things that JS was never designed to do.



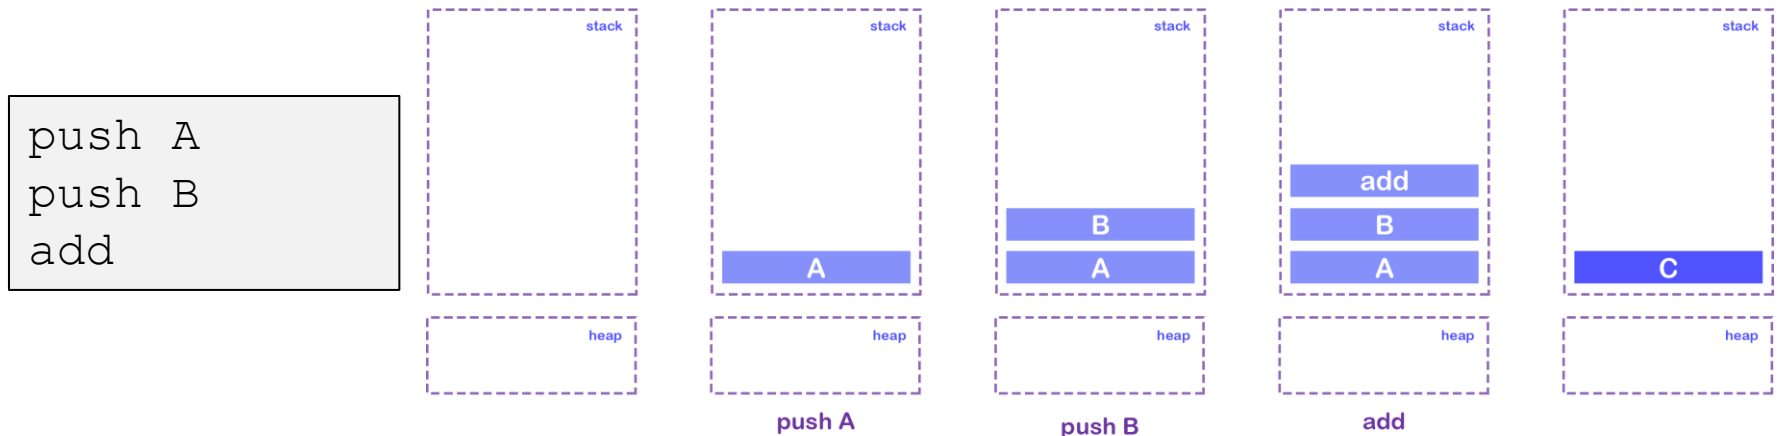
LLVM

# Main Advantages

- 1. Performance:** WebAssembly offers strong type guarantees, it gives more consistent and reliable performance than JS.
- 2. Portability:** because you can compile from other languages, you can bring open source libraries built in languages like C++.
- 3. Flexibility:** the ability to write in other languages.
  - To date, JS has been the only fully supported option.
  - Now with WebAssembly, you get more choice.

# WebAssembly is a Stack Machine

- Stack machine is a virtual machine (like a processor) that takes one instruction at a time and pushes it on the stack.
- When an operation is needed to be performed, it will pop values from the stack and compute the result.
- Example: compute the sum of two integers:



# WebAssembly is a Stack Machine

- Here's an example of a WebAssembly function that computes the quantity  **$F(x)=2x^2+1$** .

```
(func $f (param $x i32) (result i32)
    ;; stack: []
    (get_local $x) ;; stack: [x]
    (get_local $x) ;; stack: [x, x]
    (i32.mul)      ;; stack: [x*x]

    (i32.const 2)  ;; stack: [2, x*x]
    (i32.mul)      ;; stack: [2*x*x]


    (i32.const 1)  ;; stack: [1, 2*x*x]
    (i32.add))     ;; stack: [2*x*x+1]
```

# WebAssembly is a Stack Machine

- First, the syntax: this is the WebAssembly text format, which is derived from S-expressions.
- The same format commonly used for the Lisp programming.
- An S-expression is either
  - an “atom” (e.g. `get_local`, `2`, or `$x`),
  - or a list of S-expressions surrounded by parentheses (e.g. `(get_local $x)`, `(param $x i32)`).

S-expression

(\* 2 (+ 3 4))



# WebAssembly is a Stack Machine

- **MOST** instructions in WebAssembly modify the value stack in some way.

```
(i32.const 1)  
(i32.const 2)  
(i32.sub)
```



# F(x) in a register-based notation


```
(i3f:
    movl    %edi, -4(%rbp) ; *(rbp-4) = x
    shll    $1, %edi       ; edi = edi << 1 (= edi * 2)
    imull   -4(%rbp), %edi ; edi = edi * *(rbp-4)
    addl    $1, %edi       ; edi = edi + 1
    movl    %edi, %eax     ; eax = edi
    retq                               ; return eax2.const 1)
(i32.const 2)
(i32.sub)
```

- Of course, x86 still has a stack (e.g. `-4(%rbp)`), but it uses that in tandem with registers.

# WebAssembly is Build Target

- Binary format - WASM

```
get_local 0
i64.eqz
if (result i64)
    i64.const 1
else
    get_local 0
    get_local 0
    i64.const 1
    i64.sub
    call 0
    i64.mul
end
```



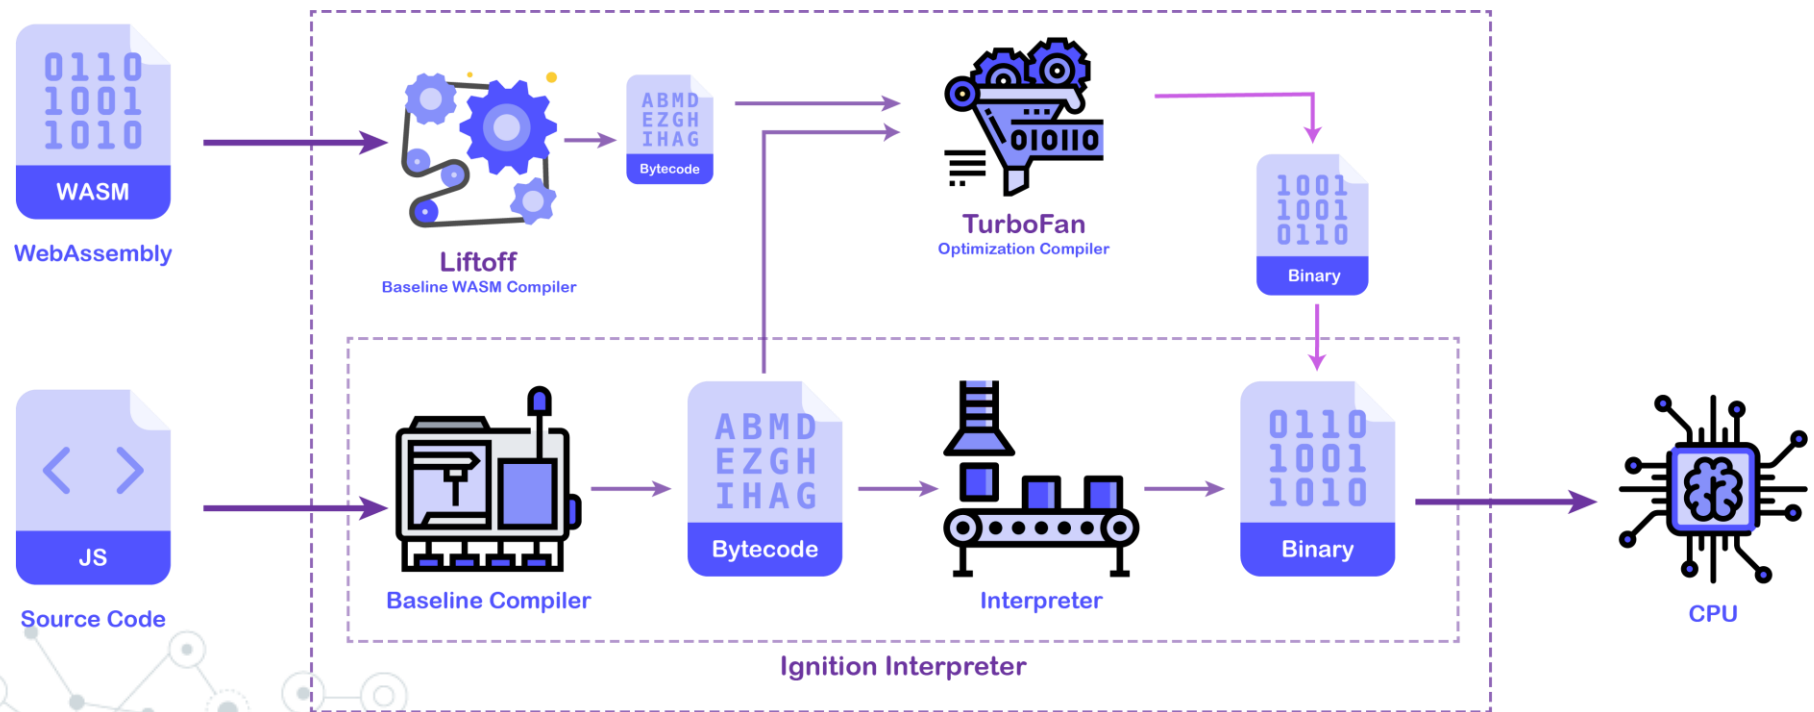
```
20 00
50
04 7E
42 01
05
20 00
20 00
42 01
7D
10 00
7E
0B
```

# How WebAssembly works?

- Every browser has a JavaScript engine which runs JavaScript.
- How it can run WebAssembly binary instructions?
- Browsers have introduced a new baseline compiler to compile WebAssembly to a bytecode that JavaScript interpreter can understand.

# How WebAssembly works?

- V8 integrated Liftoff, their WebAssembly baseline compiler whose job is to compile WebAssembly into an unoptimized bytecode as quick as possible.



# Can we use WebAssembly?

- Shipped in all major browsers.
- The first new language to ship in every major browser since JS.






# Why WebAssembly is Fast?

# Improvements to WASM

## 1- Implicit Caching (faster startup time):


- When a site loads a WebAssembly module, it first goes into the Liftoff compiler to execute it.
- Then, the code is further optimized off the main thread through the turbo fan optimizing compiler.
- Then, the results are hot swapped in when ready.



# Improvements to WASM

- Now, with implicit caching, we cache that optimized WebAssembly module directly in the HTTP cache.
- Then, after the user leaves the page and comes back, we load that optimized module directly from the cache, resulting in immediate top tier performance.


***New Style***





# Improvements to WASM

## 2- New Features: *WebAssembly Threads*



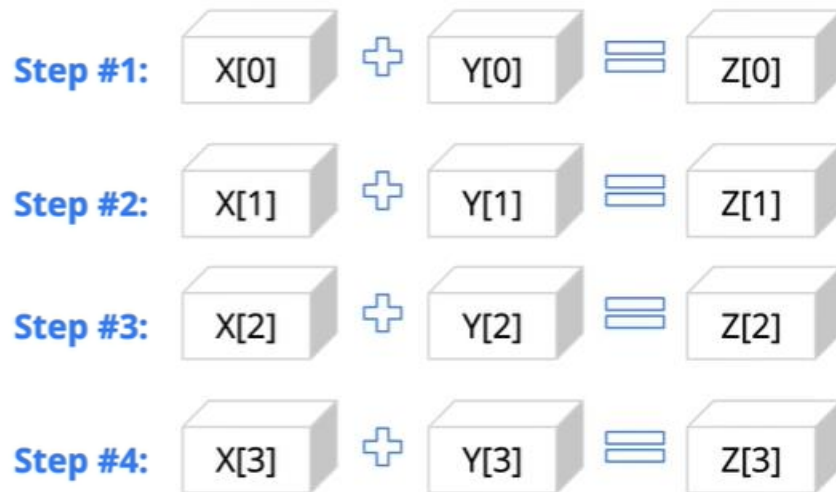
- WebWorkers: allows WebAssembly to run on different CPU cores.
- SharedArrayBuffer: allows WebAssembly to operate on the same piece of memory.
- Atomic Operations: synchronizing WebAssembly so that things happen in the right order.

# Improvements to WASM


## 2- New Features: *SIMD* (Single Instruction Multiple Data)

```
for (i = 0; i < 4; i++)  
    z[i] = x[i] + y[i]
```

### *Without SIMD*



### *With SIMD*



# Improvements to WASM

- The CPU is able to vectorize these elements and just take a single CPU operation to add them together.
- One example of SIMD is the hand tracking system:  
[Mediapipe.page.link/web](https://mediapipe.page.link/web)
- Without SIMD: we are getting about three frames per second
- With SIMD, we get a much smoother 15 frames per second.

***Please start the slide show to run the demo***



# Rust & WebAssembly Demo

- <https://m1el.github.io/wasm-asteroids/demo/demo.html>

## Rust, JavaScript and WebAssembly

```
use wasm_bindgen::prelude::*;
#[wasm_bindgen]
extern {
    pub fn alert(s: &str);
}
#[wasm_bindgen]
pub fn greet(name: &str) {
    alert(&format!("Hello, {}!", name));
}
```



```
use wasm_bindgen::prelude::*;
```

- ◎ The first line contains a use command, which imports code from a library into your code.
- .
- ◎ wasm-bindgen provides a bridge between the types of JavaScript and Rust. It allows JavaScript to call a Rust API with a string, ....

## Calling external functions in JavaScript from Rust

```
#[wasm_bindgen]
extern {
    pub fn alert(s: &str);
}
```

- ◎ extern, which tells Rust that we want to call some externally defined functions.
- ◎ The attribute says "wasm-bindgen knows how to find these functions".
- ◎ The third line is a function signature, written in Rust. It says "the alert function takes one argument, a string named s."
- ◎ The alert function provided by JavaScript. We call this function in the next section.

## Producing Rust functions that JavaScript can call

```
#[wasm_bindgen]
pub fn greet(name: &str)
{
    alert(&format!("Hello,
{}!", name));
}
```

Once again, we see the `#[wasm_bindgen]` attribute. In this case, it's not modifying an extern block, but a `fn`; this means that we want this Rust function to be able to be called by JavaScript. It's the opposite of `extern`: these aren't the functions we need, but rather the functions we're giving out to the world.



## The code in detail ....

- ◎ This function is named `greet`, and takes one argument, a string (written `&str`), `name`. It then calls the `alert` function we asked for in the `extern` block above. It passes a call to the `format!` macro, which lets us concatenate strings.
- ◎ The `format!` macro takes two arguments in this case, a format string, and a variable to put in it. The format string is the `"Hello, {}!"` bit. It contains `{}`s, where variables will be interpolated. The variable we're passing is `name`, the argument to the function, so if we call `greet("Steve")` we should see `"Hello, Steve!"`.
- ◎ This is passed to `alert()`, so when we call this function, we will see an alert box with `"Hello, Steve!"` in it.